



Off the Vine

An Oracle 19c & MySQL Project

By Charlie Evert
DSS 640

Introduction

Background

Off the Vine is a fictional company that sells Tomatoes and Strawberries worldwide for several months every year. As an internationally sought-after brand, Off the Vine must harness the power of data analysis to set KPI objectives with. It has several data sources, including its sales journal, its product catalog, its location catalog (where its products are sold), its customer list (sourced from each transaction) and its payment processing services.

The following write-up will summarize ways that Off the Vine can normalize its enterprise's data and integrate it into Oracle 19c so that it may effectively capitalize upon the power of data driven insights.

Scope

The data sources that Off the Vine uses have many common entities, all of which can be easily normalized through the use of Primary and Foreign keys. Off the Vine was careful to name tables and columns with the same value with the same name, as well, which simplifies the normalization process. The goal of this project is to consolidate the data from each of its sources into a cohesive data warehouse that can seamlessly integrate into Off the Vine's analytical needs.

Off the Vine's full process of data normalization and integration will be covered in this report, from its schema reconciliation to its star schema (with pertinent metadata included). Several queries, groupings, views and applications will also be explored in this report.

Scope limitations include access to data sources (the majority of the data was simulated), the absence of a live connection to transaction data and the inability for data to be seamlessly integrated with Oracle 19c (instead of needing to be converted from an Excel file).

Logical Design

Individual Schemas

Sales Journal

SALES(Trans_ID, Prod_ID, Location_ID, Empl_ID, Cust_ID, Payment_Type, Price, Qty, Total)

TRANSACTIONDATE(Trans_ID, Date)

Product Catalog

PRODUCT(Prod_ID, Prod_Name, Prod_Type)

PRESICIDE(Pest_Manu, Prod_ID)

Location Catalog

LOCATION(Location_ID, Neighborhood, State/Province, Country)

Customer List

CUSTOMER(Customer_ID, First_Name, Last_Name, DOB)

Payment Processing

Payment_Type(Payment_Type, Decription)

Reconciled Schema

Due to the fact that this data is fictionally simulated, and was ideally constructed, no aspects need extensive reconciliation. Each separate table, taken from Excel, can be seamlessly reconciled by way of foreign keys (as demonstrated below, they are referred to by _FK). Additionally, I designated primary keys with the suffix _PK.

Off the Vine Enterprise Data Warehouse

SALES(Trans_ID_PK, Prod_ID, Location_ID, Empl_ID, Cust_ID, Payment_Type, Price, Qty, Total)

TRANSACTIONDATE(Trans_ID_PK_FK, Date)

PRODUCT(Prod_ID_PK_FK, Prod_Name, Prod_Type)

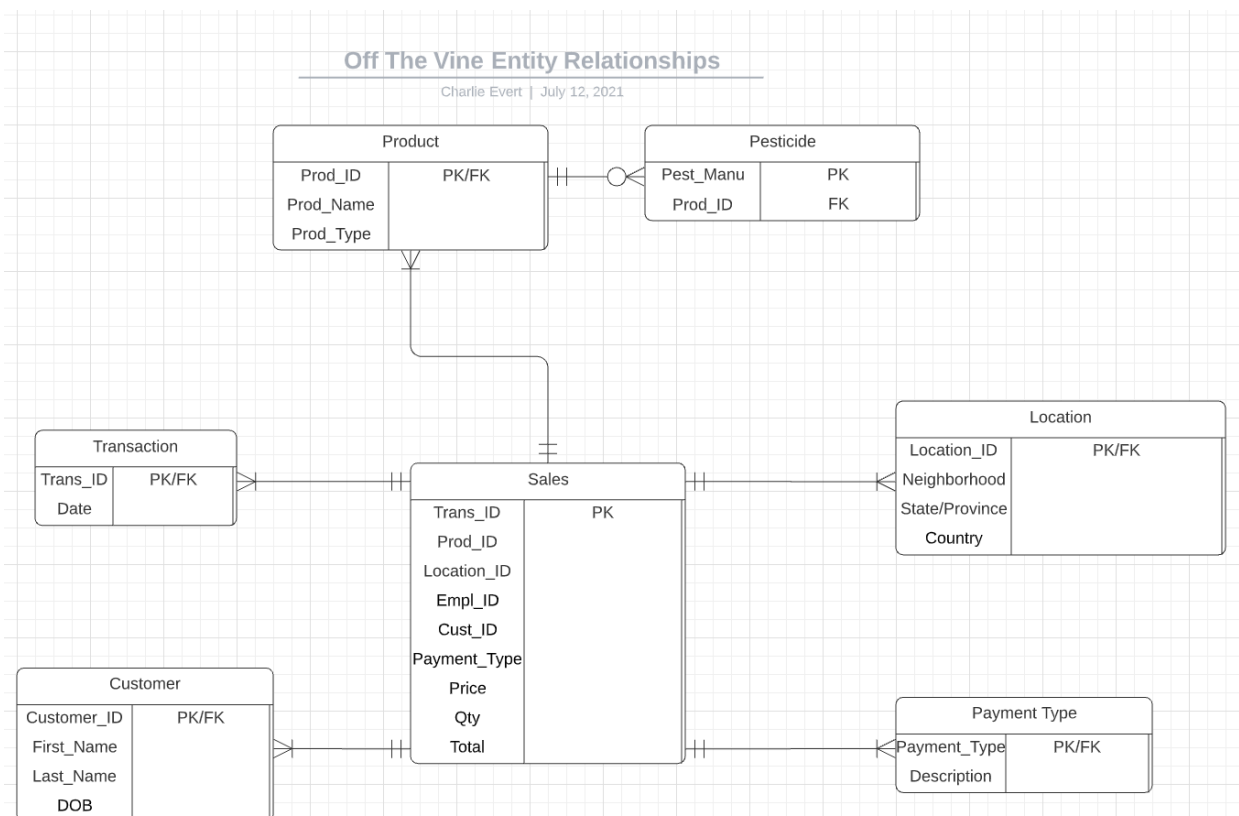
PRESICIDE(Pest_Manu_PK, Prod_ID_FK)

LOCATION(Location_ID PK FK, Neighborhood, State/Province, Country)

CUSTOMER(Customer_ID PK FK, First_Name, Last_Name, DOB)

Payment_Type(Payment_Type PK FK, Decription)

Pertinent Meta Data & Final Conceptual/Fact Schemata



The above Entity Relationship Diagram illustrates the final star/fact schema used for Off the Vine enterprise's database. Most relationships illustrated are one to many relationships between the central table in the star schema (sales) and its branched tables; this is due to the fact that each table is interconnected, and helps to describe some aspect of a simulated ledger.

Please see Appendix A for constraints.

The following is pertinent metadata:

- Trans_ID, Prod_ID, Location_ID, Empl_ID, Cust_ID, & Qty are all of the Integer (INT) data type.
- Prod_Name, Prod_Type, First_Name, Last_Name, Neighborhood, State/Province, Country & Description are all String (VARCHAR) data types.
- DOB & Date are both Date and Time (DATETIME) data types.

- Price and Total are both Float (NUMERIC) data types.

Data Staging

To start, tables were created via the CREATE TABLE function with the data types listed above (as Pertinent Meta Data).

Then, data sourced from Excel was imported into Oracle 19c via the 'Actions...' button in Oracle 19c.

This data did not need any transformation due to it being randomized and simulated appropriately; there were no missing values to worry about. Values were simulated using various RANDBETWEEN functions in Excel.

Implementation

Aggregate Queries

The following aggregate queries help to explain the data via aggregations. The number of transactions, as well as the maximum/minimum amount of money spent in one transaction were analyzed. There is a wide range of values, ranging from \$.50 to \$20 over the course of 500 transactions. See appendix B for outputs for each aggregate query.

```
SELECT COUNT(Trans_ID) AS "Number of Transactions"  
FROM SALES;
```

The picture above illustrates a COUNT query, which returns N. N is useful for determining how many sales it took to run out of inventory (or simply the number of sales between periods). This can help to set benchmarks for subsequent periods.

```
SELECT MAX(Total) AS "Maximum Money Spent in One Transaction"  
FROM SALES;
```

The picture above illustrates a MAX query, which returns the maximum total value of a transaction. This helps identify which customers to incentivize for repeat visits (so that they spend as much as they can).

```
SELECT MIN(Total) AS "Minimum Money Spent in One Transaction"  
FROM SALES;
```

The picture above illustrates a MIN query, which returns the minimum total value of a transaction. This is helpful in order for analytics professionals to help incentivize lower paying customers to buy more, via upselling and promotions.

Join Queries

The following join queries join multiple tables and display key disaggregated values. Each provides a handy joining of attributes between tables that can display data derived from key business operations. Outputs are visible in Appendix C.

```
SELECT VINECUSTOMER.First_Name AS First, VINECUSTOMER.Last_Name AS Last, SALES.Total
FROM SALES
INNER JOIN VINECUSTOMER ON SALES.Cust_ID=VINECUSTOMER.Customer_ID;
```

The above query would help managers at Off the Vine to identify customers that spend money multiple times on their products. These customers are valuable, and data indicating who they are is helpful for future advertising.

```
SELECT DISTINCT(LOCATION.Neighborhood), LOCATION.State_Province AS "State or Province", LOCATION.Country, SALES.Total
FROM LOCATION
INNER JOIN SALES ON LOCATION.Location_ID=SALES.Location_ID;
```

This query demonstrates each neighborhood's distinct total sales volumes. This is helpful in determining which towns would be best to export to (since towns with higher sales volume have the potential to garner more sales volume through increased product availability). It also lists the country of the town, which is helpful in determining exports as well since it makes sense to export more tomatoes and strawberries to countries have higher demand for them (in this case, Canada's demand is superior to the United States).

Groupings

The following Grouping queries are helpful to Off the Vine's analysts in identifying the relative performance of key attributes. This is helpful to managers in that ill-performing or bottlenecking operations identified via grouping can be trimmed or discontinued to increase the bottom line. Appendix D details the outputs of the following groupings.

```
SELECT COUNT(Cust_ID), Payment_Type
FROM SALES
GROUP BY Payment_Type
ORDER BY COUNT(Cust_ID) DESC;
```

The above grouping demonstrates a breakdown of how many customers used each payment type. In this case, few people used 2nd payment method (card) whereas significantly more used checks. One could conclude that, for whatever reason, processes must be optimized for the facilitating of payment via checks for Off the Vine's strawberries and tomatoes; perhaps local customs are responsible and the majority Canadian consumer enjoys the act of check writing. Managers should address this, given the bottlenecking associated with spending a few minutes writing a check instead of swiping a card.

```

SELECT
    COUNT(sales.cust_id),
    sales.location_id,
    system.location.neighborhood,
    system.location.country
FROM
    sales
    INNER JOIN system.location ON sales.location_id = system.location.location_id
GROUP BY
    sales.location_id,
    system.location.neighborhood,
    system.location.country
ORDER BY
    COUNT(sales.cust_id) DESC

```

The grouping query above demonstrates the count of customers that bought from Off the Vine per neighborhood. This is similar to the second join query in the previous section, but is superior given its ease of use. This query is, like its predecessor, useful for identifying target markets as well as underperforming markets. Managers should use this data to adjust exports to each neighborhood accordingly.

Materialized Views

Views are helpful for creating repeatable, easy to read outputs that, if used with a dynamic data set, can be up to date after every time a view is selected. The following views provide broad overviews of attributes between tables, and each view can be integral towards managerial decision-making. Appendices E demonstrate outputs.

Create or Replace View neighborhood AS _

```

SELECT
    system.product.prod_name,
    system.location.neighborhood,
    system.location.country,
    system.vinecustomer.first_name,
    system.vinecustomer.last_name
FROM
    system.location
    INNER JOIN system.sales ON system.sales.location_id = system.location.location_id
    INNER JOIN system.product ON system.sales.prod_id = system.product.prod_id
    INNER JOIN system.vinecustomer ON system.sales.cust_id = system.vinecustomer.customer_id;

SELECT * FROM Sales_Journal
ORDER BY neighborhood DESC;

```

The above view, “neighborhood”, displays each sale broken down into which product was sold, where, and to whom. This is helpful in determining repeat customers and what they bought by volume, so that exports can be adequate to maximize sales.


```
CREATE OR REPLACE VIEW Above_Average_Totals AS
SELECT Trans_ID, Total
FROM SALES
WHERE Total > (SELECT AVG(Total) FROM SALES);

SELECT * FROM Above_Average_Totals
```

Above_Average_Totals view, pictured above, demonstrates a repeatable view that calls all transaction IDs (along with their associated totals) that were above average. Above average transactions may illuminate trends regarding particular time frames that lead to higher-than-average totals; the variables creating such periods could then be further analyzed to determine how to incentivize customers to have larger totals.

```
CREATE OR REPLACE VIEW STDEV_TOTAL AS
SELECT STDDEV(Total) AS "Standard Deviation"
FROM SALES;


SELECT * FROM STDEV_TOTAL;
```

The STDEV_TOTAL view above shows the overall standard deviation of totals (per transaction). This standard deviation of approximately \$6 works well with the Above_Average_Totals view above to demonstrate exactly how rare a given total would be. Being able to call this function rapidly helps analysts to see if standard deviations (and variations) are changing given new data. If standard deviations suddenly change, this is either a cause of concern or evidence that a marketing decision to spur increase sales per transaction is working.

Business Analytics Applications

Artificially Intelligent Forecasts via JMP

Artificial intelligence can help to predict revenues for each location or employee, among other predictive analytics. After dummy coding Cash, Card and Bitcoin and forming multivariate models that correlate with revenues (per transaction), it is clear that there is little difference amongst payment methods in predicting revenues. Please see below for JMP's artificially intelligent output (which is reflective of random inputs being used).



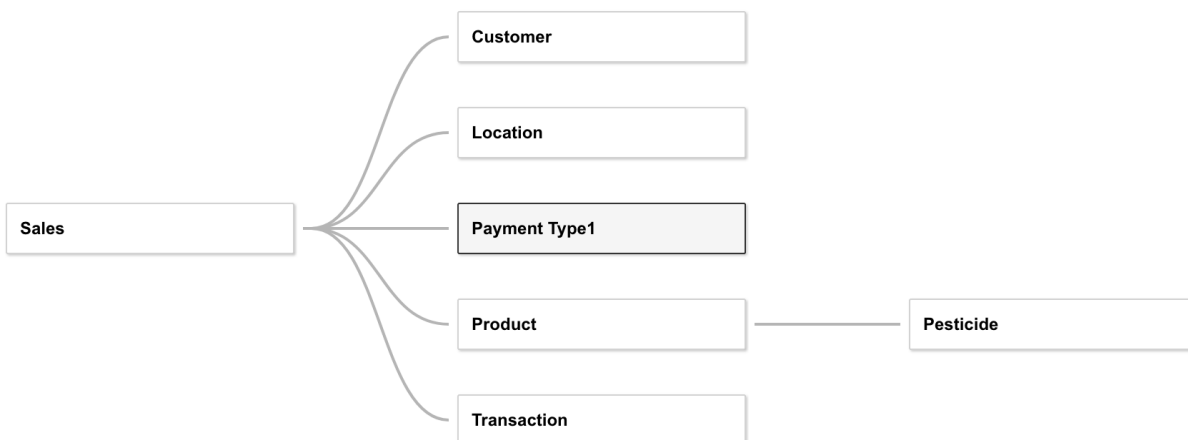
Below is JMP's prediction estimate for revenue given payment types. From this equation, we can conclude that Bitcoin payments typically lead to the highest amount of sales revenues whereas card payments lead to the lowest amount of sales revenues.

▼ Prediction Expression

$$1.347826087 + -0.097826087 \cdot \text{Cash} + -0.163615561 \cdot \text{Card} + -0.074016563 \cdot \text{Bitcoin}$$

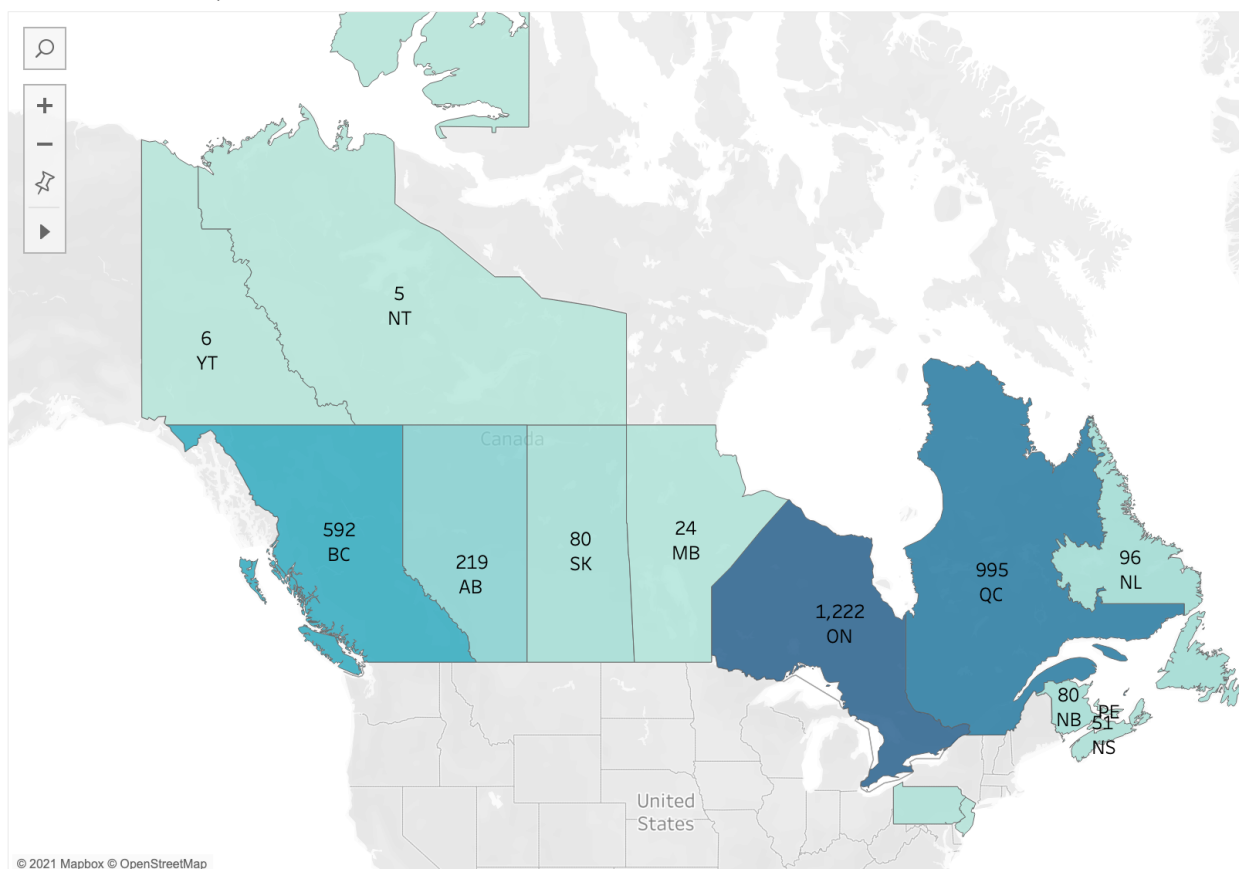
Niche Identification via Tableau

IoT can help to identify customer trends, including finding particular locations with people more likely to purchase from Off the Vine, helping further distribution and the profitability of advertising. IoT applies sales revenue numbers since cash registers and payment processing systems automatically upload this information online, making the data easily accessible to Off the Vine. Please see below for how easily a star schema is assembled in Tableau for the purpose of analysis given clean, normalized enterprise data.



Once data is assembled into a star schema, attributes that aren't connected in a single table (but are connected in a data warehouse by foreign keys) can lead to easily identifying trends. Using Geographic Information Systems in Tableau derived from Off the Vine's data, I constructed a heatmap viewable below that demonstrates a high concentration of total sales as being on the Canadian coasts, with very few in the United States.

Sales Heatmap



From this application, it can be gleamed that Coastal Canadians really love Off the Vine's strawberries and tomatoes; analytics, here, demonstrates that advertising to regions that aren't Quebec, British Columbia or Ontario is not important to revenues. Thus, these other regions should not be sold to if advertising costs are equivalent in each region. As an aside, these findings are likely due to the fact that most neighborhoods gathered for the purpose of random sampling were in these regions; the data itself follows a normal distribution reminiscent of simulation.

Conclusion

This project was great in bridging gaps between various data managerial and analytical tools that I have frequently worked with. I learned that it can be quite simple to build complete databases at the enterprise level through the use of Oracle and MySQL.

Some short comings of this project include:

- Data was simulated, and therefore insights from this data are difficult to come by (other than illustrating the central limit theorem)
- Data was static and unchanging, so changes in queries (especially views) cannot be observed and their utility is therefore lessened
- I found difficulty in regulating referential integrity constraint, so I kept the data as akin to a star schema as possible to minimize any errors

In future projects, I will:

- Source real data with the potential for actionable insights
- Work with live data sources
- Explore more obscure aggregative functions

Overall, I felt that this project's practicum helped me grasp database management far more effectively than watching lectures/videos and reading due to the hands-on nature of database management at the enterprise level.

Appendix A: Constraints

DSS 630.sqlWelcome PageDSS 630LOCATIONColumns | Data | Model | Constraints | Grants | Statistics | Triggers | Flashback | D

Actions...

| | CONSTRAINT_NAME | CONSTRAINT_TYPE | SEARCH_CONDITION |
|---|-----------------|-----------------|------------------|
| 1 | LOCATION_PK | Primary Key | (null) |

DSS 630.sqlWelcome PageDSS 630PRODUCTColumns | Data | Model | Constraints | Grants | Statistics | Triggers | Flashback | D

Actions...

| | CONSTRAINT_NAME | CONSTRAINT_TYPE | SEARCH_CONDITION |
|---|-----------------|-----------------|------------------|
| 1 | PRODUCT_PK | Primary Key | (null) |

DSS 630.sqlWelcome PageDSS 630PESTICIDEColumns | Data | Model | Constraints | Grants | Statistics | Triggers | Flashback | D

Actions...

| | CONSTRAINT_NAME | CONSTRAINT_TYPE | SEARCH_CONDITION |
|---|-----------------|-----------------|------------------|
| 1 | PESTICIDE_PK | Primary Key | (null) |
| 2 | PEST_FK | Foreign Key | (null) |

DSS 630.sqlWelcome PageDSS 630PAYMENT_TYPEColumns | Data | Model | Constraints | Grants | Statistics | Triggers | Flashback | Dependencies |

Actions...

| | CONSTRAINT_NAME | CONSTRAINT_TYPE | SEARCH_CONDITION | R_OWNER |
|---|-----------------|-----------------|------------------|---------|
| 1 | PAYMENT_TYPE_PK | Primary Key | (null) | (null) |

DSS 630.sqlWelcome PageDSS 630VINECUSTOMERColumns | Data | Model | Constraints | Grants | Statistics | Triggers | Flashback | Dependencies |

Actions...

| | CONSTRAINT_NAME | CONSTRAINT_TYPE | SEARCH_CONDITION | R_OWNER |
|---|-----------------|-----------------|------------------|---------|
| 1 | CUSTOMER_PK | Primary Key | (null) | (null) |

DSS 630.sqlWelcome PageDSS 630TRANSACTION_DATEColumns | Data | Model | Constraints | Grants | Statistics | Triggers | Flashback | Dependencies |

Actions...

| | CONSTRAINT_NAME | CONSTRAINT_TYPE | SEARCH_CONDITION | R_OWNER |
|---|-----------------|-----------------|------------------|---------|
| 1 | TRANS_DATE_PK | Primary Key | (null) | (null) |

DSS 630.sqlWelcome PageDSS 630SALESColumns | Data | Model | Constraints | Grants | Statistics | Triggers | Flashback | Dependencies | Details | Partitions | Indexes | SQL

Actions...

| | CONSTRAINT_NAME | CONSTRAINT_TYPE | SEARCH_CONDITION | R_OWNER | R_TABLE_NAME | R_CONSTRAINT_NAME |
|---|-----------------|-----------------|------------------|---------|------------------|-------------------|
| 1 | CUST_FK | Foreign Key | (null) | SYSTEM | VINECUSTOMER | CUSTOMER_PK |
| 2 | LOCATION_FK | Foreign Key | (null) | SYSTEM | LOCATION | LOCATION_PK |
| 3 | PMT_TYPE_FK | Foreign Key | (null) | SYSTEM | PAYMENT_TYPE | PAYMENT_TYPE_PK |
| 4 | PROD_FK | Foreign Key | (null) | SYSTEM | PRODUCT | PRODUCT_PK |
| 5 | SALES_PK | Primary Key | (null) | (null) | (null) | (null) |
| 6 | TRANS_DATE_FK | Foreign Key | (null) | SYSTEM | TRANSACTION_DATE | TRANS_DATE_PK |

Appendix B: Aggregate Queries

The screenshot shows the DSS 630 Query Builder interface. The top toolbar includes icons for running queries, saving, and other functions. The 'Worksheet' tab is active, displaying the following SQL query:

```
SELECT COUNT(Trans_ID) AS "Number of Transactions"
FROM SALES;
```

Below the query editor, the 'Script Output' tab shows the execution results:

```
Number of Transactions
-----
500
```

The status bar indicates 'Task completed in 0.031 seconds'.

The screenshot shows the DSS 630 Query Builder interface. The top toolbar includes icons for running queries, saving, and other functions. The 'Worksheet' tab is active, displaying the following SQL query:

```
SELECT MAX(Total) AS "Maximum Money Spent in One Transaction"
FROM SALES;
```

Below the query editor, the 'Script Output' tab shows the execution results:

```
Maximum Money Spent in One Transaction
-----
20
```

The status bar indicates 'Task completed in 0.03 seconds'.

The screenshot shows the DSS 630 Query Builder interface. The top toolbar includes icons for running queries, saving, and other functions. The 'Worksheet' tab is active, displaying the following SQL query:

```
SELECT MIN(Total) AS "Minimum Money Spent in One Transaction"
FROM SALES;
```

Below the query editor, the 'Script Output' tab shows the execution results:

```
Minimum Money Spent in One Transaction
-----
.5
```

The status bar indicates 'Task completed in 0.025 seconds'.

Appendix C: Join Queries

DSS 630.sql Welcome Page DSS 630 VINECUSTOMER 0.177 seconds

Worksheet Query Builder

```
SELECT VINECUSTOMER.First_Name AS First, VINECUSTOMER.Last_Name AS Last, SALES.Total
FROM SALES
INNER JOIN VINECUSTOMER ON SALES.Cust_ID=VINECUSTOMER.Customer_ID;
```

Script Output Task completed in 0.177 seconds

| FIRST | LAST | TOTAL |
|----------|-----------|-------|
| Jonah | Lomu | 10 |
| Antoni | Nichols | 20 |
| Tammy | Felipa | 10 |
| Cole | Murphy | 14 |
| Nieves | Gordon | 6 |
| Austin | Bryant | 12 |
| Rosamond | Epifania | 18 |
| Hugo | Stevens | 4 |
| Dylan | Gutierrez | 2.5 |

DSS 630.sql Welcome Page DSS 630 LOCATION 0.145 seconds

Worksheet Query Builder

```
SELECT DISTINCT (LOCATION.Neighborhood), LOCATION.State_Province AS "State or Province", LOCATION.Country, SALES.Total
FROM LOCATION
INNER JOIN SALES ON LOCATION.Location_ID=SALES.Location_ID;
```

Script Output Task completed in 0.145 seconds

| Milton | ON CA | 8 |
|------------------|-------|-------|
| Thunder Bay | ON CA | 4 |
| Red Deer | AB CA | 4 |
| Sault Ste. Marie | ON CA | 4 |
| NEIGHBORHOOD | St CO | TOTAL |
| Wood Buffalo | AB CA | 3 |
| New Westminster | BC CA | 16 |
| Saint-Jérôme | QC CA | 18 |
| Granby | QC CA | 4 |
| St. Thomas | ON CA | 12 |
| Port Coquitlam | BC CA | 4 |
| North Bay | ON CA | 8 |
| Brandon | MB CA | 16 |
| Rimouski | QC CA | 2 |
| Georgina | ON CA | 1.5 |
| Vernon | BC CA | 5 |
| NEIGHBORHOOD | St CO | TOTAL |

Appendix D1: Groupings

The screenshot shows a SQL query editor window with a toolbar at the top. The toolbar includes icons for running the query, saving, and other standard SQL IDE functions. The text '0.066 seconds' is displayed on the right side of the toolbar. Below the toolbar, there are two tabs: 'Worksheet' and 'Query Builder'. The 'Worksheet' tab is active, showing a SQL query:

```
SELECT COUNT(Cust_ID), Payment_Type  
FROM SALES  
GROUP BY Payment_Type  
ORDER BY COUNT(Cust_ID) DESC;
```

Below the query editor, there is a 'Script Output' window. It shows the result of the query execution, which is a table with two columns: 'COUNT(CUST_ID)' and 'PAYMENT_TYPE'. The table contains four rows of data:

| COUNT(CUST_ID) | PAYMENT_TYPE |
|----------------|--------------|
| 138 | 4 |
| 126 | 3 |
| 122 | 1 |
| 114 | 2 |

Appendix D2: Groupings

DSS 630.sql x Welcome Page x DSS 630 x SALES x

0.097 seconds

Worksheet Query Builder

```

SELECT
  COUNT(sales.cust_id),
  sales.location_id,
  system.location.neighborhood,
  system.location.country
FROM
  sales
  INNER JOIN system.location ON sales.location_id = system.location.location_id
GROUP BY
  sales.location_id,
  system.location.neighborhood,
  system.location.country
ORDER BY
  COUNT(sales.cust_id) DESC

```

Script Output x

Task completed in 0.097 seconds

| COUNT(SALES.CUST_ID) | LOCATION_ID | NEIGHBORHOOD | CO |
|----------------------|-------------|----------------|----|
| 12 | 212 | Squamish | CA |
| 12 | 111 | Rimouski | CA |
| 10 | 99 | Port Coquitlam | CA |
| 10 | 179 | La Prairie | CA |
| 10 | 158 | North Cowichan | CA |
| 10 | 47 | Cambridge | CA |
| 10 | 135 | Prince Albert | CA |
| 10 | 199 | Essex | CA |
| 10 | 18 | Kitchener | CA |
| 8 | 58 | Waterloo | CA |
| 8 | 220 | Thorold | CA |

| COUNT(SALES.CUST_ID) | LOCATION_ID | NEIGHBORHOOD | CO |
|----------------------|-------------|--------------|----|
| 8 | 9 | Calgary | CA |
| 8 | 39 | Sherbrooke | CA |
| 8 | 61 | Red Deer | CA |

Appendix E1: Materialized Views

DSS 630.sql x Welcome Page x DSS 630 x SALES x

0.211 seconds

Worksheet Query Builder

```

SELECT
    system.product.prod_name,
    system.location.neighborhood,
    system.location.country,
    system.vinecustomer.first_name,
    system.vinecustomer.last_name
FROM
    system.location
    INNER JOIN system.sales ON system.sales.location_id = system.location.location_id
    INNER JOIN system.product ON system.sales.prod_id = system.product.prod_id
    INNER JOIN system.vinecustomer ON system.sales.cust_id = system.vinecustomer.customer_id;

SELECT * FROM Sales_Journal
ORDER BY neighborhood DESC;

```

Script Output x

Task completed in 0.211 seconds

| | | | |
|------------|--------------|--------------|----------------------|
| Tomato | Beaconsfield | CA Luca | Turner |
| Tomato | Beaconsfield | CA Minda | Cody |
| PROD_NAME | NEIGHBORHOOD | CO | FIRST_NAME LAST_NAME |
| Tomato | Beaconsfield | CA Minda | Cody |
| Tomato | Beaconsfield | CA Frederick | Gardner |
| Tomato | Beaconsfield | CA Frederick | Gardner |
| Tomato | Beaconsfield | CA Michael | Reed |
| Tomato | Beaconsfield | CA Michael | Reed |
| Strawberry | Bathurst | CA Theo | Thomas |
| Strawberry | Bathurst | CA Theo | Thomas |
| Tomato | Bathurst | CA Pamala | Iesha |
| Tomato | Bathurst | CA Pamala | Iesha |
| Strawberry | Barrie | CA Kristi | Dane |
| Strawberry | Barrie | CA Quinn | Carpenter |
| PROD_NAME | NEIGHBORHOOD | CO | FIRST_NAME LAST_NAME |
| Strawberry | Barrie | CA Kristi | Dane |
| Strawberry | Barrie | CA Quinn | Carpenter |
| Tomato | Baie-Comeau | CA Spencer | Armstrong |
| Tomato | Baie-Comeau | CA Spencer | Armstrong |
| Strawberry | Aurora | CA Colette | Corine |
| Tomato | Aurora | CA Colette | Corine |
| Tomato | Aurora | CA Colette | Corine |

Appendix E2: Materialized Views

DSS 630.sql x Welcome Page x DSS 630 x SALES x

0.073 seconds

Worksheet Query Builder

```
CREATE OR REPLACE VIEW Above_Average_Totals AS
SELECT Trans_ID, Total
FROM SALES
WHERE Total > (SELECT AVG(Total) FROM SALES);

SELECT * FROM Above_Average_Totals
```

Script Output x

Task completed in 0.073 seconds

| | |
|----------|-------|
| 432 | 8 |
| 434 | 20 |
| 435 | 12 |
| 442 | 8 |
| 450 | 14 |
| 451 | 16 |
| ----- | |
| TRANS_ID | TOTAL |
| ----- | |
| 454 | 18 |
| 455 | 8 |
| 456 | 8 |
| 457 | 18 |
| 459 | 20 |
| 462 | 12 |
| 464 | 10 |
| 465 | 12 |
| 471 | 10 |
| 472 | 12 |
| 477 | 18 |
| ----- | |
| TRANS_ID | TOTAL |
| ----- | |
| 478 | 16 |
| 480 | 20 |
| 487 | 10 |
| 494 | 20 |
| 495 | 12 |
| 499 | 10 |

171 rows selected.

Appendix E3: Materialized Views

The screenshot shows a SQL IDE interface with the following components:

- Top Bar:** Contains tabs for 'DSS 630.sql', 'Welcome Page', 'DSS 630', and 'SALES'. Below the tabs is a toolbar with various icons and a status bar indicating '0.081 seconds'.
- Worksheet Tab:** The 'Query Builder' sub-tab is active. It displays the following SQL script:

```
CREATE OR REPLACE VIEW STDEV_TOTAL AS
SELECT STDDEV(Total) AS "Standard Deviation"
FROM SALES;

SELECT * FROM STDEV_TOTAL;
```
- Script Output Tab:** The 'Script Output' sub-tab is active. It displays the results of the SQL execution:

```
View STDEV_TOTAL created.

Standard Deviation
-----
5.83134531
```